

Which Loans Might Default? Lending Club Data Analysis and Prediction



Meskerem Goshime

September 22, 2022

Data Wrangling and Data Cleaning

- I started with loan data from 2007-2015 with 74 columns and 759,339 rows.
- Columns with 50k or more missing values were dropped.
- Null values were imputed with 0 or 'Other' where it made sense.
- After that, remaining rows with missing values were dropped.
- Took care of inconsistency in data entry by combining similar values in some columns.


Outliers Handling

- The most significant outliers were in annual_inc and dti columns.
- Max Annual income = 9 million, 99.7 percentile 379 k
- Max dti = 380, 99.7 percentile 39
- Rows with values beyond the 99.7 percentile in the respective columns were removed.


Feature Selection

- Columns with redundant data were dropped.
- Dropped columns that seem unnecessary based on my intuition from the data exploration
- I ended up with 33 columns from 74.
- In the end I selected 10 columns out of the 33 using Feature Importance.

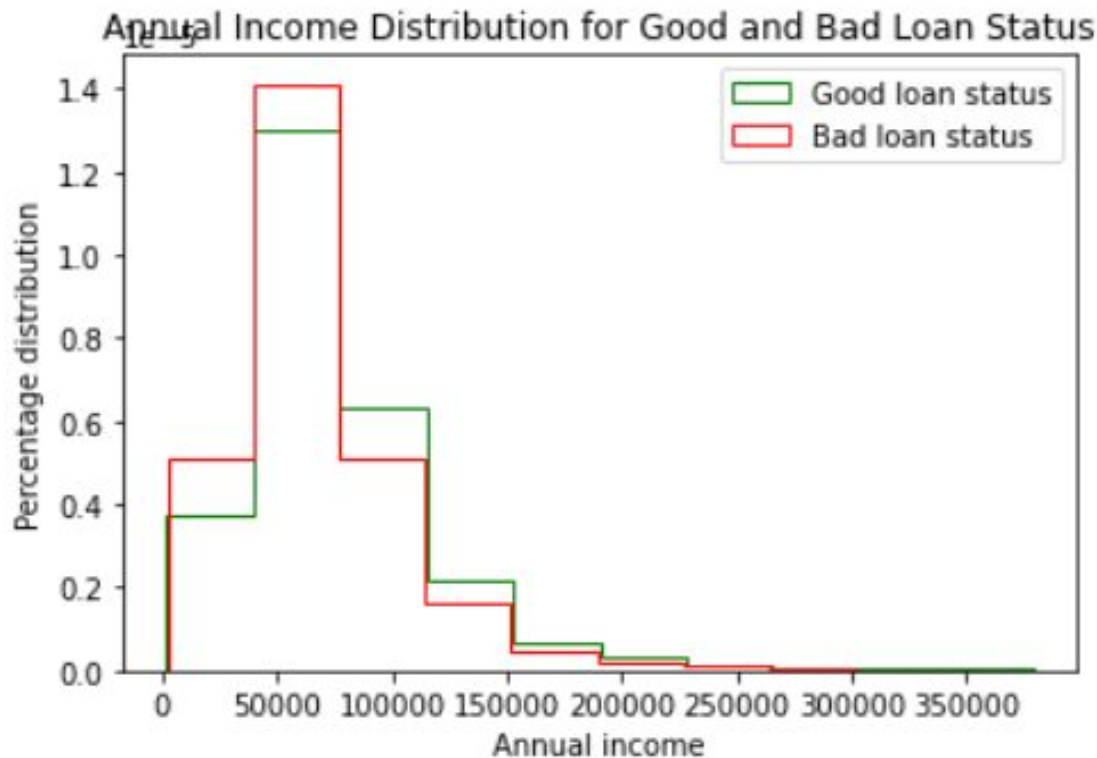
EDA - Interest Rate of Borrowers in Bad Loan Status is Significantly Higher



EDA - Bad Loan Statuses Correspond with Higher Median DTI Value



EDA - Lower DTI Corresponds with Higher Annual Income



EDA - Zip Codes with Highest Default Ratio


Zip Code	Loans in Bad Status	Loan Count	% in Bad Status
415xx	12	75	0.160
736xx	12	83	0.144
237xx	26	191	0.136
126xx	28	209	0.133
638xx	20	154	0.129
668xx	13	105	0.123

Preprocessing & Training Data Development

- Grade and sub-grade columns were encoded as numeric columns.
- The Categorical columns were encoded using One Hot Encoding.
- The numeric columns were standardized using StandardScaler.
- Target and Predictor Variables.
 - ◆ Loan status was chosen as the target variable (y).
 - ◆ The rest of the columns became the predictor variables (X).
- The data was split into Training Set, X_train, y_train (80%) and Test Set, X_test, y_test (20%).


Modeling - Random Forest Classifier

- Chosen n-estimator value: 500
- ROC/AUC Score: 0.74




Modeling - Logistic Regression

→ ROC/AUC Score: 0.72



Modeling - Gradient Boosting

- n_estimator value: 600
- Max_depth: 3
- ROC AUC Score: 0.77



Best Performing Model - Gradient Boosting

Best performing model - Gradient Boosting

- Best ROC/AUC score 0.77
- Minimizes the false positives while also keeping the false negatives low.

Planned Improvements

- Try PCA dimensionality reduction to see if the performance of the model will improve.
- Try resampling method to handle the imbalance in the data.

**“All models are wrong, but some models
are useful.”**

George E. P. Box

Project Files

Project Notebooks

Project Report

Special Thanks

- For Husain Battiwala for making [the data available on Kaggle!](#)
- For Tony Paek for his amazing mentorship!
- For my husband and boys for their encouragement and support!